

NOTES FOR MEETING

Starting too narrow, don't know what I'm applying machine learning techniques to

Structure of problem is not well understood

deep learning vs reinforcement learning

state vs. stateless model

DIFFERENT KINDS OF NEURAL NETWORKS

INTRODUCTION TO REINFORCEMENT LEARNING

3.1 K-ARMED BANDIT PROBLEM

We would like to maximize expected reward based on k different actions, assuming you do not know the true rewards. For each action A_t that takes a value a , we then have $q_*(a) = E(R_t | A_t = a)$. We maintain an estimate of $q_*(a)$, $Q_t(a)$. We can either take a "greedy" action by choosing what we know to be the highest, or "explore" by choosing a different action. Thus in the short run, we do better by being greedy, but in the long run potentially better by exploring, so what actions we take depends on how many time steps remain.

3.2 ACTION-VALUE METHODS

The simplest way to estimate $Q_t(a)$ is to replace the population expectation with sample averages:

$$Q_t(a) = \frac{\sum_{i=1}^{t-1} R_i 1(A_i = a)}{\sum_{i=1}^{t-1} 1(A_i = a)}$$

which is consistent for $q_*(a)$ by the law of large numbers. In this sense, the greedy action A_t can be described by $\operatorname{argmax}_a Q_t(a)$. Further, instead of always being greedy, we can choose to explore with probability ε and use an ε -greedy method. The advantage is that asymptotically, every action is sampled, with little harm to optimality. Positive values of ε almost always outperform pure greediness, and this is especially true for nonstationary rewards.

When it comes to estimating the value of an action, we can use the formula earlier to get

$$Q_n = \frac{R_1 + R_2 + \dots + R_{n-1}}{n-1}$$

However, it is more computationally efficient to use an updating rule instead of holding records for every selection:

$$NewEstimate \leftarrow OldEstimate + StepSize[Target - OldEstimate]$$

Under the previous estimator, we would have $Q_{n+1} = Q_n + \frac{1}{n}[R_n - Q_n]$, and Q_1 is an arbitrary selection.

3.3 NONSTATIONARY PROBLEM AND INITIAL VALUES

One way R changing over time is to provide more weight to recent rewards, which can be done by using a constant step-size parameter $\alpha \in (0, 1]$. Our updating formula then becomes $Q_{n+1} = (1 - \alpha)^n Q_1 + \sum_{i=1}^n \alpha(1 - \alpha)^{n-i} R_i$.

It is clear that with constant step size, the bias induced by the initial choice for Q_1 does not go away. This has advantages and disadvantages; on the one hand, it can be used to encourage exploration early on (the method has more leeway to choose non-greedy actions in the beginning), but at the expense of making that assumption. We refer to this as an optimistic initial value, and it is a simple but effective way to explore with stationary rewards.

3.4 UPPER-CONFIDENCE-BOUND AND GRADIENT BANDIT ALGORITHMS

Instead of randomly exploring ε often, we can take into account how fruitful that action may be. If we choose actions according to:

$$A_t = \operatorname{argmax}_a \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$

Further, we may be able to train the algorithm to have a preference over its actions; i.e. for each action value a , there is a preference $H_t(a)$. Then the probability of taking that action, based on the softmax function, is:

$$P(A_t = a) = \frac{e^{H_t(a)}}{\sum_{b=1}^k e^{H_t(b)}} = \pi_t(a)$$

The learning/updating method for these preferences is based on Stochastic Gradient Ascent¹:

$$\begin{aligned} H_{t+1}(A_t) &= H_t(A_t) + \alpha(R_t - \overline{R_t})(1 - \pi_t(A_t)) \\ H_{t+a}(a) &= H_t(a) - \alpha(R_t - \overline{R_t}) \end{aligned}$$

Demeaning the rewards serves to check if the latest reward is high or low compared to baseline, and if it is high, then the preference is positively updated, and vice versa.

POLICIES AND MARKOV DECISION PROCESSES

In our framework, there is an agent and an environment, and at each time step, the agent makes a decision by selecting an action. Then, the environment is updated and changed, and the agent selects another action. The means by which the agent gains information about the environment is a representation, known as a state $S_t \in \mathcal{S}$; the consequent action is $A_t \in \mathcal{A}$ with a reward $R_t \in \mathcal{R}$. A_t, S_t , and R_t are random variables that take on where $\mathcal{A}, \mathcal{S}, \mathcal{R}$ are the set of all actions and states.

The agent seeks to maximize expected return G_t , some function of the rewards. It is nice to use a discounted sum of the form:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}, \gamma \in [0, 1]$$

A policy $\pi : \mathcal{S} \rightarrow \mathbb{R}$ is a mapping from a state to a probability of selecting an action, where if the agent is following the policy π , $P(A_t = a | S_t = s) = \pi(a|s)$. The state-value function v_π is the expected return in the state s of following π :

$$v_\pi(s) = E(G_t | S_t = s)$$

and the action-value function q_π :

$$q_\pi(s, a) = E(G_t | S_t = s, A_t = a)$$

The value functions can be estimated by sample averages of the rewards received under certain

¹It can be shown that the expected value of the update is equal to the expected value of the gradient of the reward